Comprehensive Report: Fundamentals of Generative AI and Large Language Models

Name: Shrivarshan PG

Register no: 212225240146

Slot: T1-D6

Subject: 19DT608 – Prompt Engineering

1.1 Subject

Generative AI, LLMs, and 2024–2025 AI Tool Landscape

1.2 Target Audience

Students and Technical Professionals

1.3 Table of Contents

---

1. Executive Summary

This report provides a foundational overview of Generative AI, tracing its evolution from early machine learning to modern Large Language Models (LLMs). It explains model training, highlights key AI tools in 2024, and discusses the Transformer architecture that powers modern AI systems.

---

2. Introduction to Generative AI

Artificial Intelligence (AI) has evolved from rule-based systems to models capable of generating human-like content such as text, images, audio, and videos. Generative AI and LLMs are central to this transformation.

Generative AI focuses on creating new content rather than just analyzing data.

2.1 Foundational Concepts

- **Generative Model:** A model that learns how data is generated to create new data
- **Probability Distribution:** Mathematical framework for predicting likelihoods

---

Introduction to Artificial Intelligence and Machine Learning

- **Artificial Intelligence (AI):** Simulation of human intelligence in machines
- **Machine Learning (ML):** Subset of AI that learns patterns from data

Types of Machine Learning

- Supervised Learning
- Unsupervised Learning
- Reinforcement Learning

---

3. What is Generative AI?

Generative AI creates new content such as text, images, music, and code.

Key Characteristics

- Learns data distributions
- Produces realistic content
- Uses deep learning and probability

Examples: ChatGPT, DALL·E, MusicLM

3.2 Types of Generative Models

2.2 Types of Generative Models

Model Type	Mechanism	Primary Use Case
GANs (Generative Adversarial Networks)	Two networks (Generator and Discriminator) compete. One creates, the other critiques.	High-quality image generation, deepfakes.
VAEs (Variational Autoencoders)	Encodes input into a compressed space and decodes it back to create variations.	Image denoising, molecular discovery.
Diffusion Models	Gradually adds noise to data and then learns to reverse the process to "recover" an image.	Midjourney, DALL-E 3, Stable Diffusion.

3.3 Generative Models and Their Types:

A Generative Model learns the underlying distribution of data and generates new samples similar to the training data.

---

4. Generative Models and Their Types

A generative model learns patterns in data and produces similar outputs.

4.1 Generative Adversarial Networks (GANs)

- Two networks: Generator & Discriminator
- Used in image generation and deepfakes

4.2 Variational Autoencoders (VAEs)

- Encode data into latent space
- Decode to generate new data

4.3 Diffusion Models

- Remove noise gradually to generate data
- Used in Stable Diffusion and DALL·E

---

5. The 2024 AI Toolscape

Popular Tools

Text & Reasoning:

- GPT-4o (OpenAI)
- Claude 3.5 Sonnet (Anthropic)
- Gemini 1.5 Pro (Google)

Image Generation:

- Midjourney v6
- DALL·E 3
- Adobe Firefly

Video Generation:

- Sora (OpenAI)
- Runway Gen-3
- Luma Dream Machine

Coding Tools:

- GitHub Copilot
- Cursor
- Replit Ghostwriter

5.1 2024 AI Tools Some popular AI tools used in 2024 include:

Tool Name	Purpose
ChatGPT	Text generation, tutoring
Google Gemini	Multimodal AI
Claude	Long-document reasoning
Midjourney	Image generation
Stable Diffusion	Open-source image generation
GitHub Copilot	Code assistance
Canva AI	Design and content creation

---

6. Large Language Models (LLMs): How They Are Built

An LLM is a neural network trained to predict the next word in a sequence.

6.1 Building Blocks

- Transformer Architecture
- Self-Attention Mechanism
- Positional Encoding

6.2 Training Process

- **Pre-training:** Learning from large datasets
- **Supervised Fine-Tuning (SFT):** Learning from examples
- **RLHF:** Aligning with human feedback

Capabilities of LLMs

- Answer questions
- Summarize text
- Translate languages
- Generate code

Examples: GPT-3, GPT-4, BERT, PaLM

---

7. Architecture of Large Language Models

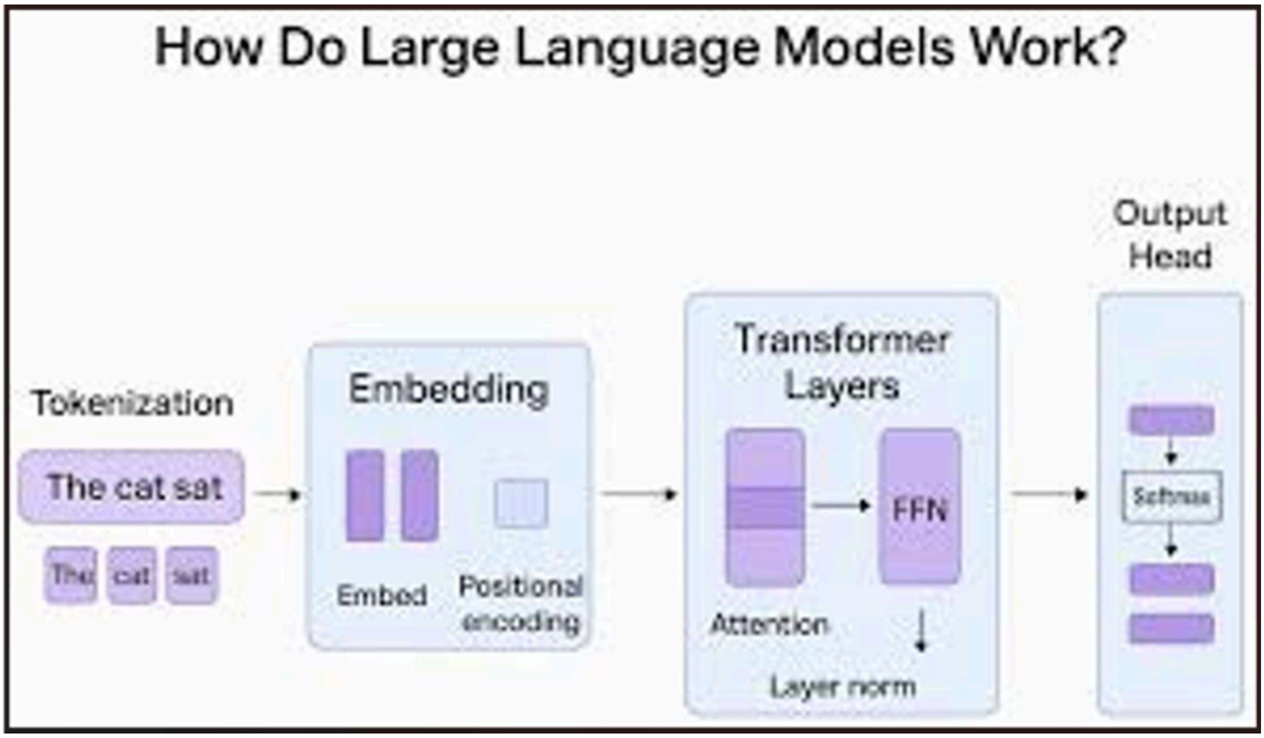

7.1 Transformer Architecture

Key components:

- Tokenization
- Embedding
- Self-Attention
- Feedforward Networks
- Output Layer

7.2 Popular LLM Architectures

- **GPT (Generative Pre-trained Transformer) – Decoder-only model**
- **BERT (Bidirectional Encoder Representations from Transformers) – Encoderonly model**

---

8. Training Process and Data Requirements

Training Steps

- Data collection
- Data cleaning
- Pre-training
- Fine-tuning
- Deployment

Data Requirements

- Large-scale text data
- High-quality datasets
- Powerful hardware (GPUs/TPUs)

---

9. Use Cases and Applications

Major Applications:

- Chatbots and assistants
- Content creation
- Code generation
- Healthcare
- Education
- Business automation

---

10. Technical Implementation (Example Code)

Next-Token Prediction Logic

LLMs function essentially as advanced probability machines. Below is a conceptual implementation of how "Self-Attention" weights might be calculated.

```
import numpy as np

def calculate_attention(query, key, value):
```

 # Step 1: Compute attention scores (Dot Product)

```
scores = np.dot(query, key.T)
```

 # Step 2: Normalize with Softmax (sum of weights = 1)

```
weights = np.exp(scores) / np.sum(np.exp(scores))
```

 # Step 3: Weighted sum of values

```
output = np.dot(weights, value)
return output, weights


# Example: Finding context for the word "Bank" (River bank vs Money bank)
q = np.array([1, 0, 0]) #"Query" for river context
k = np.array([1, 0, 0]) #"Key" for the word "Water"
v = np.array([10, 20, 30]) #"Value" containing semantic data

context_result, attention_weights = calculate_attention(q, k, v)
print(context_result)
```

---

11. Limitations and Ethical Considerations

Limitations

- High computational cost
- Data bias
- Hallucinations

- Limited real-world understanding

Ethical Issues

- Privacy concerns
- Misinformation
- Job displacement
- Copyright issues

---

12. Future Trends

- Multimodal AI
- Smaller, efficient models
- Personalized assistants
- Better safety and alignment
- Expansion in healthcare and education

---

13. Evolution of AI (Timeline)

The following chart traces the journey from early logic to

modern generative power:

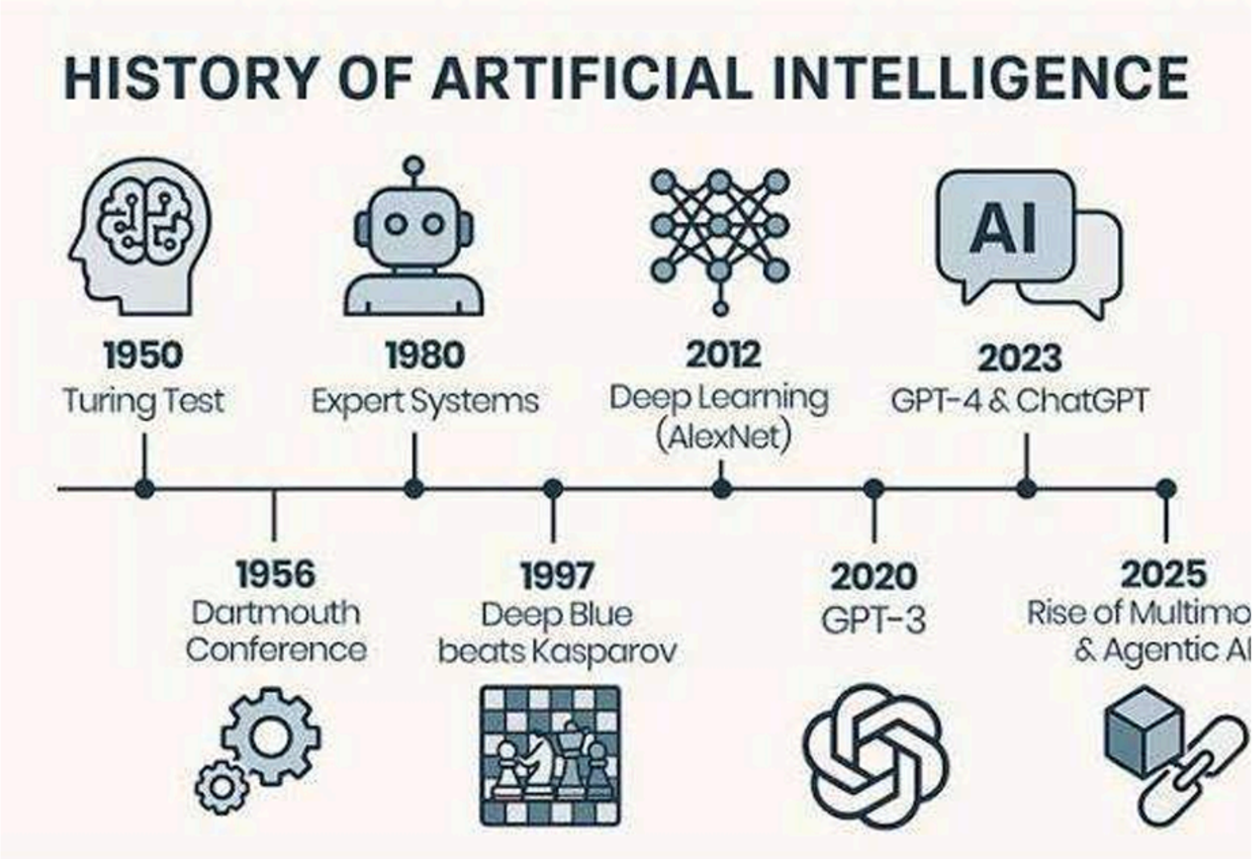

- **1950: Alan Turing proposes the "Turing Test."**
  1956: Dartmouth Conference; the term "Artificial Intelligence" is coined.
  1966: ELIZA, the first chatbot, is created at MIT.
  1997: IBM's Deep Blue defeats world chess champion Garry Kasparov.
  2012: AlexNet triggers the "Deep Learning" revolution.
  2017: Google publishes "Attention is All You Need" (The birth of Transformers).
  2022: OpenAI launches ChatGPT, bringing Generative AI to the mainstream.
  2024–2026: Rise of Multimodal "Omni" models and Autonomous AI Agents.

---

14. Conclusion

Generative AI and LLMs represent a major advancement in AI. Understanding these technologies helps students prepare for AI-driven careers. Despite challenges, responsible development can unlock massive potential.

References

- OpenAI Documentation
- Google AI Blogs
- Deep Learning (Ian Goodfellow)
- Transformer research papers
- AI educational resources